

•••

# Reinforcement Learning and Autonomous Systems (CS4122)



Lecture 19 (28/09/2024)  
Lecture 20 (30/09/2024)

Instructor: Gourav Saha

# Lecture Content



(This lecture is the beginning of Module 3)

- Learning in MDPs.
  - Need for learning.
  - Fundamental idea behind learning in MDPs.
  - Generalized Policy Iteration.
- Monte Carlo Policy Evaluation.
  - Estimating the value function.
- Temporal Difference (TD) Policy Evaluation.
  - Estimating the value function.

# Learning in MDPs

Recall that an MDP is defined by:

- States and the associate state space.
- Actions and the associated actions space.
- Reward.
- State transition probability.
- Reward probability.

# Learning in MDPs : Need for Learning

Recall that an MDP is defined by:

- States and the associate state space.
- Actions and the associated actions space.
- Reward.
- State transition probability.
- Reward probability.

The need for learning in MDPs arises in two different scenarios:

- **Scenario 1: Involved probability distributions are not known.** For MDPs, state transition probability and reward probability are the involved probability distribution. Example:
  - Logistic networks (Amazon).
  - Scheduling in wireless networks.

In reality, some amount of knowledge of the involved probability distributions are known.

# Learning in MDPs : Need for Learning

Recall that an MDP is defined by:

- States and the associate state space.
- Actions and the associated actions space.
- Reward.
- State transition probability.
- Reward probability.

The need for learning in MDPs arises in two different scenarios:

- Scenario 2: The models are known but they are too complex to the point that using any tools like value/policy iteration is impractical. The complexity arises because:
  - State space is too large (this increase the time complexity of value/policy iteration). Example: Atari games.
  - The setup does not explicitly have an “MDP friendly” model. These setups, though MDPs, are better understood in terms of ODEs (or stochastic ODEs). Example: Bicycles, quadcopters.

# Learning in MDPs : Need for Learning

Recall that an MDP is defined by:

- States and the associate state space.
- Actions and the associated actions space.
- Reward.
- State transition probability.
- Reward probability.

The need for learning in MDPs arises in two different scenarios:

- Scenario 1: Involved probability distributions are not known.
- Scenario 2: The models are known but they are too complex to the point that using any tools like value/policy iteration is impractical.

For scenario 2, we can learn the optimal policy by simulating the model in computer and using this simulation as a proxy of real-world.

For scenario 1, we still need some model to simulate in computer. Otherwise, direct real-world deployment can be risky.

# Learning in MDPs : Need for Learning

Given: A threshold,  $\theta$ .

(S1): For all  $x \in \mathcal{S}$ , initialize  $V(x)$  arbitrarily to any real value. Initialize  $\Delta = \infty$ .

(S2): while  $\Delta > \theta$ :

(S3): for all  $x \in \mathcal{S}$ :

$$(S4): \quad V_{new}(x) = \max_{a \in \mathcal{A}(x)} \left( r(x, a) + \beta \sum_{x' \in \mathcal{S}} P[x' | x, a] V(x') \right)$$

(S5): Compute  $\Delta = \max_{x \in \mathcal{S}} |V_{k+1}(x) - V_k(x)|$ .

(S6): Set  $V(x) = V_{new}(x)$  for all  $x \in \mathcal{S}$ .

(S7): For all  $x \in \mathcal{S}$ :

$$\pi^*(x) = \operatorname{argmax}_{a \in \mathcal{A}(x)} \left( r(x, a) + \beta \sum_{x' \in \mathcal{S}} P[x' | x, a] V(x') \right)$$

(S8): Return  $\pi^*$ .

- This is the pseudocode for value iteration.
- When  $r(x, a)$  and  $P[x' | x, a]$  are not known, we can't use value iteration. This explains scenario 1 discussed in the previous slides.
- When the model is complex, i.e. too many states, this for loop which in turn is inside this while loop becomes the computational bottleneck. This explains scenario 1 discussed in the previous slides.

# Learning in MDPs : Need for Learning

Given: A threshold,  $\theta$ .

(S1): For all  $x \in \mathcal{S}$  arbitrarily initialize: (i) a policy  $\pi(x)$  to any value in  $\mathcal{A}(x)$ , and (ii) a value function  $V(x)$  to any real value. Also, set  $converged = False$ .

(S2): while not( $converged$ ):

(S3): Set Use IPE to compute the value function  $V^\pi(x)$  for all  $x \in \mathcal{S}$  corresponding to current policy  $\pi$ . Set the initial value function of IPE as  $V(x)$  and the conversion threshold as  $\theta$ .

(S4): For all  $x \in \mathcal{S}$ :

(S5): Compute,

$$\tilde{a} = \operatorname{argmax}_{a \in \mathcal{A}(x)} \left( r(x, a) + \beta \sum_{x' \in \mathcal{S}} P[x' | x, a] V^\pi(x') \right)$$

$$q^\pi(x, \tilde{a}) = r(x, \tilde{a}) + \beta \sum_{x' \in \mathcal{S}} P[x' | x, \tilde{a}] V^\pi(x')$$

➤ This is the pseudocode for policy iteration.

➤ Similar explanation as value iteration discussed in the previous slides.

(S6): If  $q^\pi(x, \tilde{a}) > V^\pi(x)$ , then set  $\pi_{new}(x) = \tilde{a}$ . Else set  $\pi_{new}(x) = \pi(x)$ .

(S7): If  $\pi_{new}(x) = \pi(x)$  for all  $x \in \mathcal{S}$ , set  $converged = True$ .

(S8): Set  $\pi(x) = \pi_{new}(x)$  and  $V(x) = V^\pi(x)$  for all  $x \in \mathcal{S}$ .

(S9): Return  $\pi$ .

# Learning in MDPs: Fundamental idea behind learning in MDPs

$$\pi_{new}(x) = \operatorname{argmax}_{a \in \mathcal{A}(x)} q^\pi(x, a)$$

- Policy iteration shows that if we can compute the Q-function,  $q^\pi(x, a)$ , corresponding to the current policy,  $\pi$ , then we can use the above equation to improve the policy.
- In module 2, we had  $r(x, a)$  and  $P[x' | x, a]$  and hence we could use value and policy iteration. But we can't do that in a learning setup.
- Another way to compute the Q-function is to realize the Q-function  $q(x, a)$  is just the expected value of the return  $G_t$  starting from state  $x$  and action  $a$ . **We can estimate the Q-function by finding the sample average of  $G_t$ .** Provide that the initial state is  $x$  and initial action is  $a$ .

This is the fundamental idea of learning in MDPs.

# Learning in MDPs: Generalized PI

- Generalized policy iteration (GPI) for the foundation of learning in MDPs.

Generalized Policy Iteration (**not rigorous** but a broadly observed trend):

- $q^\pi(x, a)$  does not have to be computed exactly (few rounds of IPE; maybe even one round which is VI).
- $q^\pi(x, a)$  does not have to be updated for all  $(x, a)$  pairs.
- Policy update:

$$\pi_{new}(x) = \operatorname{argmax}_{a \in \mathcal{A}(x)} q^\pi(x, a)$$

does not have to happen for all  $x$ .

# Learning in MDPs: Generalized PI

- Generalized policy iteration (GPI) for the foundation of learning in MDPs.

Generalized Policy Iteration (**not rigorous** but a broadly observed trend):

- Till the time policy evaluation to compute  $q^\pi(x, a)$  keeps happening for all  $(x, a)$  pairs infinitely often, AND
- Till the time policy update:

$$\pi_{new}(x) = \operatorname{argmax}_{a \in \mathcal{A}(x)} q^\pi(x, a)$$

keeps happening for all  $x$  infinitely often, THEN

- **Convergence to optimal policy is guaranteed.**

# MC Policy Evaluation

To compute the value function  $V^\pi(x)$

- The main objective is to compute the optimal policy which in turn relies on computing the Q-function for a given policy.
- But to make it simpler, we will start by computing the value function for a given policy just to get an overall idea. Remember: Computing value function for a given policy is called policy evaluation.

# MC Policy Evaluation: Trajectory

➤ Whenever we are in a “learning” setup, there is always going to be a concept of “sampling” from the environment. The question is, what to sample?

- We have to sample a **trajectory**. A trajectory is denoted using  $\tau$ . A trajectory is a **sequence** of **state, action, and reward pairs** obtained in an episode.

$$(x_0, a_0, r_0), (x_1, a_1, r_1), (x_2, a_2, r_2), \dots, (x_T, a_T, r_T)$$
$$\xleftarrow{\textcolor{blue}{t = 0}} \quad \xleftarrow{\textcolor{blue}{t = 1}} \quad \xleftarrow{\textcolor{blue}{t = 2}} \quad \xleftarrow{\textcolor{blue}{t = T}}$$

- In the trajectory shown above,  $T$  is the last time slot of the episode. Since a trajectory is a sequence, its short hand notation is,

$$\{(x_t, a_t, r_t)\}_{t=0}^T$$

- A trajectory  $\tau$  is a random variable (that's why we are sampling it) which itself consists of other random variables. So essentially,  $\tau$  can be characterized by a joint distribution. This joint distribution is dependent on the policy, reward probability, and state transition probability.
- A pseudocode to collect the trajectory is shown in the next slide.

# MC Policy Evaluation: Psuedocode to Generate Trajectory

Given: A policy,  $\pi$ .

(S1): Reset the environment to get the initial state  $x_0$ . Initialize time  $t = 0$ , and an empty list  $\tau$  that will contain the trajectory for the current episode.

(S2): while episode did not end:

(S3): Use policy  $\pi$  for the current state  $x_t$  to choose action  $a_t$ .

(S4): Take action  $a_t$ . Environment will return reward  $r_t$  and transition to next state  $x_{t+1}$ .

(S5): Append the state, action, reward pair  $(x_t, a_t, r_t)$  to  $\tau$ . Set  $t = t + 1$ .

(S6): Return trajectory  $\tau$ .

# MC Policy Evaluation

To compute the value function  $V^\pi(x)$

- The notes are not complete. You have to read the book. **That said, all the psuedocodes are there in the slides.** You have to refer the books to understand these psuedocodes.
- The following are the important concepts:
  - **First-visit and every-visit Monte Carlo:** Read chapter 5. Start from the beginning and read section 5.1. You don't have to read example 5.2 (soap bubble) if you don't want to. The psuedocode of every-visit Monte Carlo is not there in the book however it is there in [this](#) slide.

# MC Policy Evaluation: Trajectory

## Fishing in Grid World

(0.5, 3.0) • 25	(1.0, 3.0) • 26	27	28	29	(3.0, 3.0) • 30
19	20	21	22	 23	24
13	14	15	16	17	18
(0.5, 1.0) • 7	(1.0, 1.0) • 8	9	10	11	(3.0, 1.0) • 12
(0.5, 0.5) • 1	(1.0, 0.5) • 2	3	4	5	(3.0, 0.5) • 6

I did an example of collecting trajectories for first visit and every visit monte carlo during lecture. You can get the lecture notes from some one else.

Trajectory 1:

Trajectory 2:

# MC Policy Evaluation: Psuedocode for First-Visit MC

Given: A policy,  $\pi$ .

- (S1): For all  $x \in \mathcal{S}$  arbitrarily initialize: (i)  $V(x)$  to any real value, and (ii)  $N(x)$  to zero.  $V(x)$  and  $N(x)$  are the value function and the number of samples corresponding to state  $x$  respectively.
- (S2): For every episode: **For step (S3), we use the psuedocode to generate a trajectory from [this slide](#).**
- (S3): Use policy  $\pi$  to generate a trajectory  $\tau$ . Let the trajectory be as follows with the last time slot as  $T$ :  
$$(x_0, a_0, r_0), (x_1, a_1, r_1), (x_2, a_2, r_2), \dots, (x_T, a_T, r_T)$$
- (S4): For all  $x \in \mathcal{S}$ , set  $visited(x)$  to *False*. Set the return corresponding to time  $t = 0$  as follows,  
$$G_0 = \sum_{t=0}^T \beta^t r_t$$
- (S5): For  $t = 0, 1, 2, \dots, T$ :
- (S6): If not( $visited(x_t)$ ):
- (S7): Update  $V(x_t)$  as follows:  
$$V(x_t) = V(x_t) + \frac{1}{N(x_t) + 1} (G_t - V(x_t))$$
- (S8): Update  $N(x_t) = N(x_t) + 1$ . Also set,  $visited(x_t)$  to True.
- (S9): Set  $G_{t+1} = (G_t - r_t)/\beta$ .

# MC Policy Evaluation: Psuedocode for Every-Visit MC

Given: A policy,  $\pi$ .

- (S1): For all  $x \in \mathcal{S}$  arbitrarily initialize: (i)  $V(x)$  to any real value, and (ii)  $N(x)$  to zero.  $V(x)$  and  $N(x)$  are the value function and the number of samples corresponding to state  $x$  respectively.
- (S2): For every episode: **For step (S3), we use the psuedocode to generate a trajectory from [this slide.c](#)**
- (S3): Use policy  $\pi$  to generate a trajectory  $\tau$ . Let the trajectory be as follows with the last time slot as  $T$ :
- $$(x_0, a_0, r_0), (x_1, a_1, r_1), (x_2, a_2, r_2), \dots, (x_T, a_T, r_T)$$
- (S4): ~~For all  $x \in \mathcal{S}$ , set  $\text{visited}(x)$  to False.~~ Set the return corresponding to time  $t = 0$  as follows,
- $$G_0 = \sum_{t=0}^T \beta^t r_t$$
- (S5): For  $t = 0, 1, 2, \dots, T$ :
- (S6): ~~If not( $\text{visited}(x_t)$ ):~~
- (S7): Update  $V(x_t)$  as follows:
- $$V(x_t) = V(x_t) + \frac{1}{N(x_t) + 1} (G_t - V(x_t))$$
- (S8): Update  $N(x_t) = N(x_t) + 1$ . ~~Also set,  $\text{visited}(x_t)$  to True.~~
- (S9): Set  $G_{t+1} = (G_t - r_t)/\beta$ .

# MC Policy Evaluation: Psuedocode for Every-Visit MC

Given: A policy,  $\pi$ .

- (S1): For all  $x \in \mathcal{S}$  arbitrarily initialize: (i)  $V(x)$  to any real value, and (ii)  $N(x)$  to zero.  $V(x)$  and  $N(x)$  are the value function and the number of samples corresponding to state  $x$  respectively.
- (S2): For every episode: **For step (S3), we use the psuedocode to generate a trajectory from [this slide](#).**
- (S3): Use policy  $\pi$  to generate a trajectory  $\tau$ . Let the trajectory be as follows with the last time slot as  $T$ :
- $$(x_0, a_0, r_0), (x_1, a_1, r_1), (x_2, a_2, r_2), \dots, (x_T, a_T, r_T)$$
- (S4): Set the return corresponding to time  $t = 0$  as follows,
- $$G_0 = \sum_{t=0}^T \beta^t r_t$$
- (S5): For  $t = 0, 1, 2, \dots, T$ :
- (S6): Update  $V(x_t)$  as follows:
- $$V(x_t) = V(x_t) + \frac{1}{N(x_t) + 1} (G_t - V(x_t))$$
- (S7): Update  $N(x_t) = N(x_t) + 1$ .
- (S8): Set  $G_{t+1} = (G_t - r_t)/\beta$ .

# MC Policy Evaluation

## Gradient Descent Viewpoint

- Recorded a small video (10 minutes) to explain this idea:

[https://youtu.be/UtsV835\\_ri4](https://youtu.be/UtsV835_ri4)

# MC Policy Evaluation: Psuedocode for Every-Visit MC

Given: A policy,  $\pi$ .

(S1): For all  $x \in \mathcal{S}$ , arbitrarily initialize  $V(x)$  to any real value.  $V(x)$  is the value function corresponding to state  $x$ .

(S2): For every episode:

(S3): Use policy  $\pi$  to generate a trajectory  $\tau$ . Let the trajectory be as follows with the last time slot as  $T$ :

$$(x_0, a_0, r_0), (x_1, a_1, r_1), (x_2, a_2, r_2), \dots, (x_T, a_T, r_T)$$

(S4): Set the return corresponding to time  $t = 0$  as follows,

$$G_0 = \sum_{t=0}^T \beta^t r_t$$

(S5): For  $t = 0, 1, 2, \dots, T$ :

(S6): Update  $V(x_t)$  as follows:

$$V(x_t) = V(x_t) + \alpha (G_t - V(x_t))$$

(S7): Set  $G_{t+1} = (G_t - r_t)/\beta$ .

$\alpha$  can vary with episode and time slot.

# Temporal Difference Policy Evaluation

- The notes are not complete. You have to read the book. **That said, all the psuedocodes are there in the slides.** You have to refer the books to understand these psuedocodes.
- The following are the important concepts:
  - **Temporal Difference (TD) Policy Evaluation:** Read chapter 6. Start from the beginning and read sections 6.1 and 6.2 completely.

# TD Policy Evaluation: Psuedocode

Given: A policy,  $\pi$ .

(S1): For all  $x \in \mathcal{S}$ , arbitrarily initialize  $V(x)$  to any real value.  $V(x)$  is the value function corresponding to state  $x$ .

(S2): For every episode:

(S3): Reset the episode. This will give the current state  $x$ .

(S4): while episode did not end:

(S5): Use policy  $\pi$  for the current state  $x$  to choose action  $a$ .

(S6): Take action  $a$ . Environment will return reward  $r$  and transition to next state  $x'$ .

(S7): Update  $V(x)$  as follows:

$$V(x) = V(x) + \alpha \left( r + \beta V(x') - V(x) \right)$$

*$\alpha$  can vary with episode and time slot.*

(S8): Set  $x = x'$ , i.e. set current state equal to next state.



Thank you